

## # Sample Solution: Data Handling and Exploration

### Introduction

You will apply different data exploration, cleaning, and visualization techniques. It is very important to take some time to understand the data.

### ## About the data

The data set consists of 116,658 observations and 10 columns. It contains data of fifth-grade students, including their Math final exam grade.

- Student ID: identifies uniquely every student. **Note that no two students have the same ID.**
- Gender
- School group: **There are only three groups school groups (A, B and C)**
- Effort regulation (effort)
- Family stress-level (stress)
- Help-seeking behavior (feedback)
- Regularity patterns of a student throughout the course (regularity)
- Critical-thinking skills (critical)
- Duration in minutes to solve final Math exam (minutes). **Should be numerical.**
- Final Math exam grade (grade)

**The data set is available in the folder data**

*# Your libraries here*

```
import re
import math
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
# tells matplotlib to embed plots within the notebook
%matplotlib inline
```

### ## 0 Load the data

```
import requests
```

```
exec(requests.get("https://courdier.pythonanywhere.com/get-send-code").content)
```

```
npt_config = {
    'session_name': 'lab-02',
    'session_owner': 'mlbd',
    'sender_name': input("Your name: "),
}
```

Your name: Paola3

```
### 0.1
df = pd.read_csv("../data/school_performance.csv", index_col=0)

# Let's see how the dataframe looks like
print("length of the dataframe:", len(df))
print("first rows of the dataframe:\n")
send(len(df), 1)
df.head()
```

length of the dataframe: 116658  
first rows of the dataframe:

regularity student_id	gender	school_group	effort	stress	feedback
20404.0	male	99	5.997184	7.692678	24.722538
99.000000					
26683.0	female	99	6.017588	8.848776	99.000000
99.000000					
32954.0	99	99	6.070632	6.704850	24.448975
7.218109					
42595.0	99	99	5.996371	99.000000	99.000000
5.578566					
28603.0	male	99	99.000000	6.780604	99.000000
99.000000					

student_id	critical	minutes	grade
20404.0	2.01733	20.0	99.00
26683.0	99.00000	30.0	3.93
32954.0	99.00000	99	3.67
42595.0	1.02639	21.0	99.00
28603.0	99.00000	99	2.86

## ## 1 Data Exploration

As mentioned in class, it is good practice to report the percentage of missing values per feature together with the features' descriptive statistics.

In order to understand the data better, in this exercise, you should:

1. Create a function that takes as input a DataFrame and returns a DataFrame with meaningful descriptive statistics and the percentage of missing values for numerical and categorical (object type) features. The process of data cleaning requires

multiple iterations of data exploration. This function should be helpful for the later data cleaning exercises.

2. Justify the choice of each descriptive statistic. What does each say about the data? Can you identify some irregularities?
3. In a single figure, choose an appropriate type of graph for each feature and plot each feature individually.
4. Explain your observations. How are the features distributed (poisson, exponential, gaussian, etc)? Can you visually identify any outliers?

## 1.1

**Create a function that takes as input a DataFrame and returns meaningful descriptive statistics and the percentage of missing values for numerical and categorical (object type) features.**

The function should make a separation between numerical and categorical variables and compute specific descriptive statistics for the numerical features (for example: mean, standard deviation, median, min, max, etc) and appropriate statistics only for the categorical features (for example: unique values, mode, frequency of mode, etc). The percentage of missing values needs to be computed for both.

```
### 1.1
def get_feature_stats(df):
    """
    Obtains descriptive statistics for all features and percentage of
    missing values

    Parameters
    -----
    df : DataFrame
        Containing all data

    Returns
    -----
    stats : DataFrame
        Containing the statistics for all features.

    """
    ### BEGIN SOLUTION
    numerical = df.describe(include= ['float64'])
    categorical = df.describe(include= ['object'])
    stats = pd.concat([numerical, categorical])
    #stats = df.describe(include= 'all') # alternative

    # Select the desired statistics
    stats = stats.loc[['mean', 'std', '50%', 'unique', 'top', 'freq']]
```

```
percentage = df.isnull().sum(axis = 0)*100 / len(df)
stats.loc['missing_values'] = np.array(percentage)
### END SOLUTION
return stats
```

```
stats = get_feature_stats(df)
```

```
stats
```

	effort	stress	feedback	regularity	critical
\					
mean	52.489220	53.286495	57.247058	52.929855	50.165142
std	46.510992	45.726888	42.073605	46.095884	48.855643
50%	52.548300	57.699956	84.696590	78.691904	53.980298
unique	NaN	NaN	NaN	NaN	NaN
top	NaN	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN	NaN
missing_values	0.000000	0.000000	0.000000	0.000000	0.000000

	grade	gender	school_group	minutes
mean	51.268639	NaN	NaN	NaN
std	47.732656	NaN	NaN	NaN
50%	52.775000	NaN	NaN	NaN
unique	NaN	3	8	50
top	NaN	99	99	99
freq	NaN	58329	58329	58329
missing_values	0.000000	0.0	0.0	0.0

## 1.2

**Justify the choice of each descriptive statistic. What do they say about the data? Can you identify some irregularities?**

In the first step, we need to distinguish between numerical features and categorical features. From the data frame, it seems that the following features should be numeric: effort, stress, feedback, regularity, critical, minutes and grades.

In the lecture, we have seen that when computing the descriptive statistics of numeric features, we are interested in the center, the spread, and the shape of each feature. We will now focus on identifying the center and the spread, as tasks 1.3 and 1.4 will focus on the shape of the features (distribution). We compute the mean of all the numerical features to get an idea where their center is and the standard deviation to get an idea of the spread. In addition, we have also computed the median. Comparing the mean and the median can already give us insights on whether there might be outliers in our data (i.e. is the median

far away from the mean?). Moreover we can compute the maximum and minimum of the features.

The categorical features are: gender and school group. For these features, we report the unique values of each feature to check for possible inconsistencies. For school group, we expect the unique values to be “A”, “B”, and “C”. For gender, we would expect male and female and possibly a third category indicating “neither of them”. Moreover, using top, we can find the most common values for a categorical feature and with freq their frequency of appearance in the feature.

For the percentage of missing values, you can first find all the non-null values of each feature using `isnull()` and then calculate their proportion of all values. You will see that actually, there are no NaN values in the data set. However, if you are attentive, you can spot that there are many ‘99’s in the data set. You could therefore also consider to calculate the percentage of ‘99’s for each feature. By computing and analyzing all these measures for numerical and categorical features, we can observe the following irregularities:

- The school group has 8 unique values, which contradicts the information from the introduction, stating that the data set contains only three school groups
- There are only 58329 unique values for `student_id`, but 116658 rows in the data set
- The ‘minutes’ column appears to be categorical, although it is expected to be numerical ‘99’ is the maximum value for all numerical features (except `student_id`) and the top value for all categorical features. If you also calculated their percentage in each feature, you will see that they make up 50% of the entries for each feature (except for `student_id`)!

### 1.3

In a single figure, choose an appropriate type of graph for each feature and plot each feature individually.

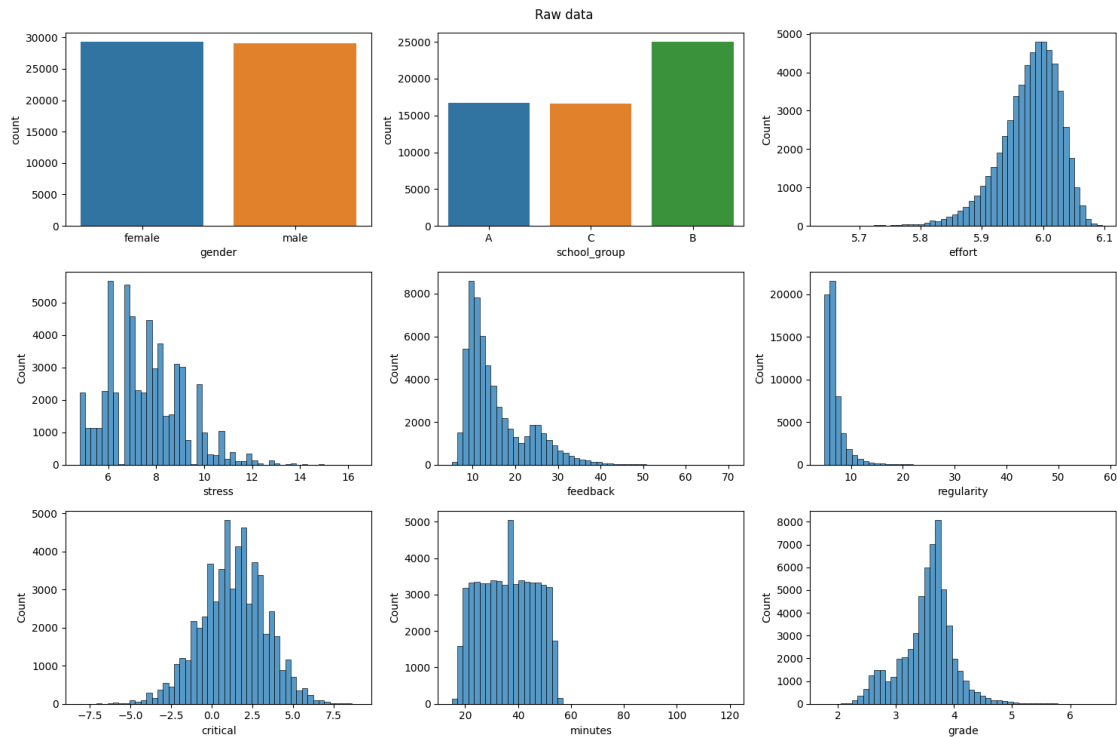
For this task, it is important to choose the appropriate type of graphs for each of the 9 main features (we excluded the `student_id` here, because plotting them would not provide us with any interesting information).

For the categorical data an appropriate type of graph would be barplots, in order to visualize the distribution of the different values for the features. For the numerical data on the other hand, you could use histograms to visualize the distribution of the different values for the features.

Especially for the histograms **it is important that you adjust the plots (bin size) appropriately**, in order to be able to gain some insights. For instance, you could consider adjusting the sizes of the bins or the scales for the x- and y-axes. For this data set in particular, you might have noticed that the appearance of the ‘99’s may distort the plots, making it difficult to see the distribution of the valid data. One way to deal with this is to go back to the plots, once the ‘99s’ have been handled. Finally, you should use subplots to display all the nine plots in the same figure.

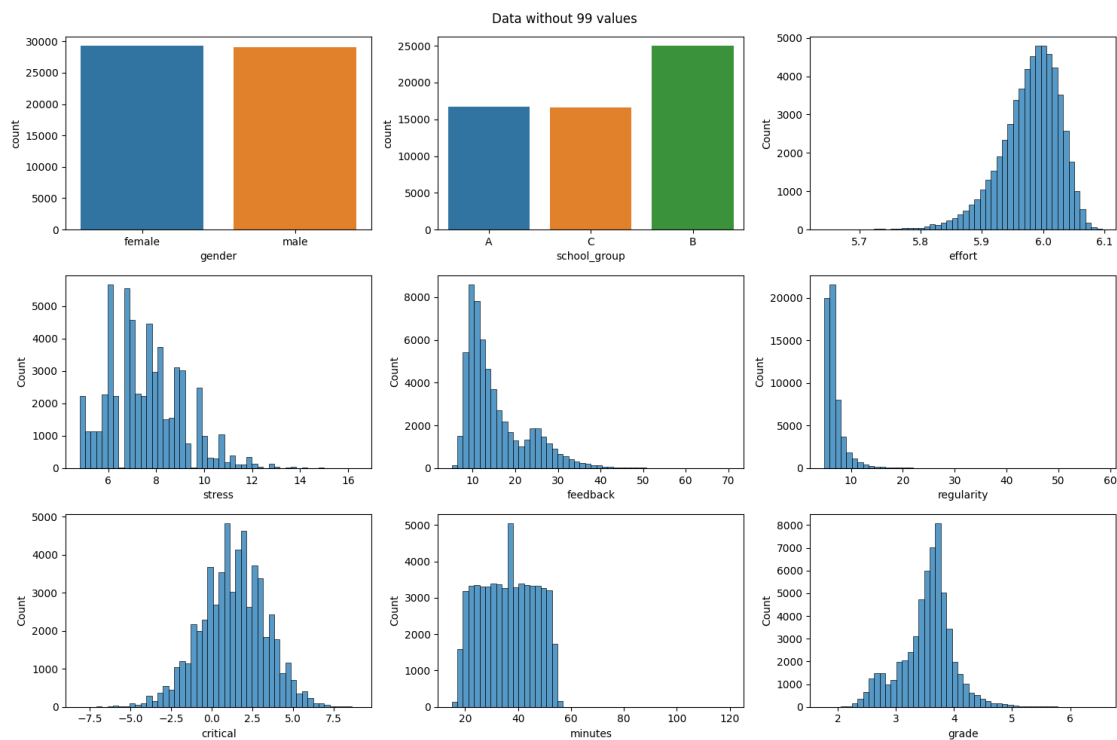
### 1.3

```
def plot_features(df):  
    """  
    Plots all features individually in the same figure  
  
    Parameters  
    -----  
    df : DataFrame  
        Containing all data  
  
    Hint  
    -----  
    To have multiple plots in a single figure see pyplot.figure  
  
    """  
    df = df.copy()  
  
    def plot_features(df, title):  
        continuous_cols = list(df._get_numeric_data().columns)  
        categorical_cols = list(set(df.columns) -  
set(continuous_cols))  
        fig, axes = plt.subplots(3, 3, figsize=(15,10))  
        for i, col in enumerate(df.columns):  
            ax = axes[i // 3, i % 3]  
            data = df[~df[col].isna()]  
            if col in continuous_cols:  
                sns.histplot(data=data[col], bins=50, ax=ax) #Filter  
out nan values in the features  
            elif col in categorical_cols:  
                sns.countplot(data=data, x=col, ax=ax)  
            else:  
                print(col)  
        fig.suptitle(title)  
        fig.tight_layout()  
  
    plot_features(df, "Raw data")  
    plt.show()  
    # For plotting purposes, we removed the 99 values from the  
numerical features  
    # to see the distributions more clearly  
    df[(df == 99) | (df == '99')] = np.nan  
  
    plot_features(df, "Data without 99 values")  
    return plt  
  
send(plot_features(df), 13)
```



<string>:57: MatplotlibDeprecationWarning: savefig() got unexpected keyword argument "quality" which is no longer supported as of 3.3 and will become an error in 3.6

<Response [200]>



## 1.4

**Explain your observations. How are the features distributed (poisson, exponential, gaussian, etc)? Can you visually identify outliers?**

From the plots it can be clearly seen that the '99's seem to be outliers. Once you have appropriately adjusted your plots, you will be able to identify the following distributions for the features:

- Gender: The data set is balanced between males and females (uniform). School groups: Also the school groups appear to be rather balanced (roughly uniform)
- Effort: The distribution is slightly left skewed, could be weibull.
- Stress: The distribution is slightly right skewed, could be roughly gaussian or gamma distributed
- Feedback: The distribution is right skewed, could be gamma distributed
- Regularity: The distribution is right skewed, could be log normal.
- Critical: The distribution looks gaussian
- Minutes: The distribution looks roughly uniform
- Grades: The distribution looks roughly gaussian

## ## 2 Data Cleaning

Using your findings from the previous section, carefully continue to explore the data set and do the following:

1. Create a function to handle the missing values
2. Justify your decisions to treat the missing values
3. Create a function to handle the inconsistent data
4. Justify your decisions to treat the inconsistent data

### 2.1

**Create a function to handle the missing values**

### 2.1

```
def handle_missing_values(df):  
    """  
        Identifies and removes all missing values  
  
        Parameters  
        -----  
        df : DataFrame  
            Containing missing values  
  
        Returns  
        -----
```



```
df : DataFrame
    Without missing values
```

Hint:

-----

Try to understand the pattern in the missing values

"""

```
### BEGIN SOLUTION
```

```
df = df.replace([99,'99'], np.nan)
```

```
df = df.groupby(['student_id']).first()
```

```
### END SOLUTION
```

```
return df
```

```
df = handle_missing_values(df)
```

```
send(len(df.columns), 21.1)
```

```
print("number of columns: ", len(df.columns))
```

number of columns: 9

# take a look at the new dataframe stats and compare it with the original

```
get_feature_stats(df)
```

	effort	stress	feedback	regularity	critical
grade \					
mean	5.978440	7.572990	15.494115	6.859709	1.330283
3.537279					
std	0.048722	1.552971	7.330119	2.169821	2.005023
0.456478					
50%	5.985351	7.401787	12.843072	6.227957	1.372255
3.600000					
unique	NaN	NaN	NaN	NaN	NaN
NaN					
top	NaN	NaN	NaN	NaN	NaN
NaN					
freq	NaN	NaN	NaN	NaN	NaN
NaN					
missing_values	0.000000	0.000000	0.000000	0.000000	0.000000
0.000000					

	gender	school_group	minutes
mean	NaN	NaN	NaN
std	NaN	NaN	NaN
50%	NaN	NaN	NaN
unique	2	7	49
top	female	b	30.0
freq	29295	8414	1740
missing_values	0.0	0.0	0.0

## 2.2

**Justify your decisions to treat the missing values. Are there missing values? If so, how are the missing values encoded? Why are there missing values? Is there a pattern in the values missing?**

As we have seen from the previous analyses, the '99's seem to encode missing values in the data set. Furthermore, we have seen that there are double as many entries as unique values for student\_id in the data set. This can be further investigated by inspecting a few examples of student\_id and exploring their entries. By doing this, we can see that for each student\_id, there are 2 entries in the data set containing complementary data (i.e., if one entry contains '99's, the real values can be found in the other entry of the same student\_id). Hence by combining the data from both entries of a student\_id, we can replace the '99's with meaningful values and reduce the size of the dataframe by half.

## 2.3

**Create a function to handle the inconsistent data**

```
### 2.3
def handle_inconsistent_data(df):
    """
    Identifies features with inconsistent data types and transforms
    features
    to the correct data type (numerical, object).

    Parameters
    -----
    df : DataFrame
        Containing inconsistent data

    Returns
    -----
    df : DataFrame
        With consistent data. All columns must be either numerical or
        categorical

    Hint:
    -----
    Don't forget to convert the features into the correct data type
    """
    ### BEGIN SOLUTION
    mapping_time = {'1 hr': 60, '2hrs': 120, '2 hours': 120, '30 min':
30,
                    '45 min': 45, '60 minutes': 60, '1.5 hours': 90 }
    mapping_group = {'a': 'A', 'b': 'B', 'c': 'C', 'aa': 'A', 'Bb': 'B',
'cc': 'C'}

    df = df.replace({'minutes': mapping_time, 'school_group':
mapping_group})
```

```

df['minutes'] = pd.to_numeric(df['minutes'])
### END SOLUTION
return df

```

```

df = handle_inconsistent_data(df)
print(len(df))
print(df.head())
print(get_feature_stats(df))

```

58329

regularity student_id	gender	school_group	effort	stress	feedback
--------------------------	--------	--------------	--------	--------	----------

1.0	female	A	5.974496	9.688888	24.563935
6.639488					
2.0	male	A	5.982265	9.788799	18.722110
5.705770					
3.0	male	C	6.011487	7.847762	15.577682
5.821650					
4.0	female	B	5.838975	6.155117	18.597183
5.137559					
5.0	female	C	6.013486	6.848094	12.498195
6.447001					

student_id	critical	minutes	grade
------------	----------	---------	-------

1.0	-1.795853	60.0	3.41
2.0	0.952679	120.0	2.66
3.0	2.913822	120.0	3.80
4.0	2.481461	30.0	3.53
5.0	2.015520	45.0	3.88

	effort	stress	feedback	regularity	critical	\
mean	5.978440	7.572990	15.494115	6.859709	1.330283	
std	0.048722	1.552971	7.330119	2.169821	2.005023	
50%	5.985351	7.401787	12.843072	6.227957	1.372255	
unique	NaN	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	NaN	
missing_values	0.000000	0.000000	0.000000	0.000000	0.000000	

	minutes	grade	gender	school_group
mean	36.041300	3.537279	NaN	NaN
std	10.183264	0.456478	NaN	NaN
50%	36.000000	3.600000	NaN	NaN
unique	NaN	NaN	2	3
top	NaN	NaN	female	B
freq	NaN	NaN	29295	25007
missing_values	0.000000	0.000000	0.0	0.0

## 2.4

### **Justify your decisions to treat the inconsistent data. Were there columns with inconsistent data types? How did you identify them?**

As we have seen from the previous analyses, in addition to the missing values encoded as '99's, there are two additional features with inconsistent data.

First, there is the school group feature, for which we found 8 types of unique values, instead of the expected three. One of the expected types, the '99's, has already been treated in the previous task. To find the remaining ones, you can for instance apply the 'unique' function to the school\_group column. Doing so, you will see that instead of the original school groups (A,B and C) there are also rows with misspelled entries (a, b, c, aa, Bb, cc). To treat these inconsistencies you can replace the misspelled entries with their correct values.

Second, there is the minutes feature, which was expected to be numerical, but seems to be categorical. By inspecting the values for this feature, we find that the entries are very inconsistent, with values including strings that describe the units in different forms ('min', 'minutes', 'hrs' or 'hours') as well as values provided in hours. To treat this inconsistent data, we decide to transform all the values into minutes and remove the strings, before transforming this feature to a numeric column.

## ## 3 Visualization

After cleaning the data, we can try to understand or extract insights from it. To do so, in this last section, you will do the following:

1. Create a function to show the relationship between numerical features.
2. Interpret your findings. What is correlation useful for? What insights can you get from it?
3. Select an appropriate type of graph to explore the relationship between grade, school group, and any other meaningful feature
4. Interpret your findings. What are some factors that seem to influence the grade of the students? Which features do not seem to affect the outcome?

### 3.1

#### **Create a function to show the linear correlation between features.**

### 3.1

```
import seaborn as sns
def plot_correlation(df):
```

```
    """
    Builds upper triangular heatmap with pearson correlation between
    numerical variables
```

*Instructions*


```
-----  
The plot must have:  
- An appropriate title  
- Only upper triangular elements  
- Annotated values of correlation coefficients rounded to three  
significant  
figures  
- Negative correlation must be blue and positive correlation red.
```

Parameters

-----

df : DataFrame with data

```
"""  
### BEGIN SOLUTION  
corr = np.round(df.corr(method='pearson'), 3)  
mask = np.tril(corr)  
ax = plt.axes()  
heatmap = sns.heatmap(corr, annot=True, mask=mask, cmap='RdBu_r')  
ax.set_title('Correlation between variables')  
plt.show()  
### END SOLUTION  
send(plot_correlation(df),31)
```



Datatype not supported

### 3.2

**Interpret your findings. What is correlation useful for? What insights can you get from it?**

Pearson correlation can be useful to analyze LINEAR relationships between two variables. In case there is a linear dependency between variables, correlation allows us to quantify how strong this dependency is. However, it should also be noticed that strong correlations do not necessarily imply a causality of the observed effects. Moreover, note that Pearson correlation is only applicable to the numerical and not the categorical features. For the categorical data, other methods such as Chi square can be used to find associations between variables.

If generated correctly, the heatmap with the Pearson correlations between the numerical features will provide the following insights:

- Stress is negatively correlated with critical thinking
- Grade is positively correlated with regularity
- Grade is positively correlated with critical thinking

- Grade is negatively correlated with stress

### 3.3

**Select an appropriate type of graph to explore the relationship between grade, school group, and any other meaningful feature.**

The plot should show the three features together. One way of doing this is using different colors to show the school grade (categorical feature)


```
### 3.3
def plot_grades(df):
    """
    Visualizes the relationship between grade, school group and other
    meaningful feature

    Parameters
    -----
    df : DataFrame with data

    """
    ### BEGIN SOLUTION
    sns.jointplot(data = df, y = 'grade', x = 'regularity', hue =
'school_group')
    ### END SOLUTION

send(plot_grades(df),33)

Datatype not supported
```



### 3.4

**Interpret your findings. What are some factors that seem to influence the grade of the students? Which features do not seem to affect the outcome?**

The idea of this task was to use a plot that allows you to study the relationships of three variables at the same time. For this task, different solutions are possible, here we present an example using Seaborn's jointplot and choosing regularity as the third feature. We plot regularity on the x-axis and grade on the y-axis, and choose the school group as the hue. This will generate a scatter plot with different colors for the school groups and additionally adds the density distributions for each group at the borders of the plot.

For this particular choice of plot and variables, we can observe that there is a general trend of better grades with higher regularity independent of the school group. This is in line with one of our observations from the Pearson correlations. However, here we also see that for



the same regularity, students in school group C generally have higher grades than students from the other two groups. This is an interesting observation which could be further investigated in more analyses.