

Lab 04- Extended Exercises on Classification and Pipelines

```
import numpy as np
import pandas as pd

from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import adjusted_mutual_info_score
from sklearn.linear_model import Lasso
from sklearn.feature_selection import SelectFromModel,
mutual_info_classif
from sklearn import model_selection
from sklearn.metrics import roc_auc_score, balanced_accuracy_score,
confusion_matrix, ConfusionMatrixDisplay, silhouette_score
from sklearn.model_selection import train_test_split
```

You are the Senior Data Scientist in a learning platform called LernTime. You have realized that many users stop using the platform and want to increase user retention. For this purpose, you decide to build a model to predict whether a student will stop using the learning platform or not.

Your data science team built a data frame in which each row contains the aggregated features per student (calculated over the first 5 weeks of interactions) and the feature dropout indicates whether the student stopped using the platform (1) or not (0) before week 10.

The dataframe is in the file `lerntime.csv` and contains the following features:

- `video_time`: total video time (in minutes)
- `num_sessions` total number of sessions
- `num_quizzes`: total number of quizzes attempts
- `reading_time`: total theory reading time
- `previous_knowledge`: standardized previous knowledge
- `browser_speed`: standardized browser speed
- `device`: whether the student logged in using a smartphone (1) or a computer (-1)
- `topics`: the topics covered by the user
- `education`: current level of education (0: middle school, 1: high school, 2: bachelor, 3: master, 4: Ph.D.).
- `dropout`: whether the student stopped using the platform (1) or not (0) before week 5.

The newest data scientist created two models with an excellent performance. As a Senior Data Scientist, you are suspicious of the results and decide to revise the code.

Your task is to:

a) Identify the mistakes. In the first cell, add a comment above each line in which you identify an error and explain the error.

b) In the second cell, you must correct the code.

```
df = pd.read_csv('data/lerntime_dropout.csv')

y = df['dropout']
X = df[['video_time', 'num_sessions', 'num_quizzes', 'reading_time',
        'previous_knowledge', 'browser_speed']]

len(df)

300
```

Task A) Identify the mistakes in the code

In the following cell, add a comment above each line in which you identify an error and explain the why it is erroneous. Please start each of your comments with `#ERROR:`. For example:

`#ERROR: the RMSE of the model is printed instead of the AUC`

```
print("The AUC of the model is: {}".format(rmse))
```

You may assume that:

- all the features are continuous and numerical.
- the features have already been cleaned and processed.

ERROR: Train-test split should be done before preprocessing steps 1. and 2. to avoid data leakage,

fitting both scaler and selector only on X_train

1. Scale the features

```
scaler = StandardScaler()
```

```
X = scaler.fit_transform(X)
```

2. Feature selection (Lasso)

```
print(X.shape)
```

```
lasso = Lasso(alpha=0.1, random_state=0).fit(X, y)
```

```
selector = SelectFromModel(lasso, prefit = True)
```

```
X = selector.transform(X)
```

```
print(X.shape)
```

3. Split the data

ERROR: The split should be done before the feature selection or transformation

to avoid data leaking

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=0)
```

Model 1

```

clf = RandomForestClassifier(n_estimators=10, max_depth=1,
random_state=0)
# ERROR: Fit should only be called on the train set
clf.fit(X,y)
preds = clf.predict(X_test)
# ERROR: The adjusted mutual information is not an appropriate score
for classification, since it would give
# a perfect score even if the predictions are the complete opposite of
y_test
print("Score model 1:
{}".format(np.round(adjusted_mutual_info_score(preds, y_test), 2)))

## Model 2
clf = RandomForestClassifier(n_estimators=1000, max_depth=None,
random_state=0)
# ERROR: Fit should only be called on the train set
clf.fit(X,y)
preds = clf.predict(X_test)
# ERROR: The adjusted mutual information is not an appropriate score
for classification, since it would give
# a perfect score even if the predictions are the complete opposite of
y_test
print("Score model 2:
{}".format(np.round(adjusted_mutual_info_score(preds, y_test), 2)))

# ERROR: The second model has just more complexity and can hence
better overfit to the test set, which leaked during training
## Discussion
# Our second model achieved perfect results with unseen data and
outperforms the first model.
## This is because we increased the number of estimators.

(300, 3)
(300, 3)
Score model 1: 0.05
Score model 2: 1.0

```

Task B) Correct the code

Correct all the erroneous code in the following cell.

```

## 1. Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.2, random_state=0)

## 2. Scale the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

## 3. Feature selection (Lasso)

```

```

print(X_train.shape)
lasso = Lasso(alpha=0.1, random_state=0).fit(X_train, y_train)
selector = SelectFromModel(lasso, prefit = True)
X_train = selector.transform(X_train)
X_test = selector.transform(X_test)
print(X_train.shape)

## Model 1
clf = RandomForestClassifier(n_estimators=10, max_depth=1,
random_state=0)
clf.fit(X_train, y_train)
preds = clf.predict(X_test)
print("Score model 1:
{}".format(np.round(balanced_accuracy_score(preds, y_test), 2)))

## Model 2
clf = RandomForestClassifier(n_estimators=1000, max_depth=None,
random_state=0)
clf.fit(X_train, y_train)
preds = clf.predict(X_test)
print("Score model 2:
{}".format(np.round(balanced_accuracy_score(preds, y_test), 2)))

## Discussion
# Our first model outperformed the second model.
# However, it is not clear why because we change the number of
estimators and the maximum depth at the same time

(240, 3)
(240, 3)
Score model 1: 0.9
Score model 2: 0.81

```

Task C) Re-write your code using pipelines.

Hint: Go over sklearn-pipeline-introduction.

```

from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV, StratifiedGroupKFold

scalers = [
    StandardScaler(),
    'passthrough'] # none

feature_selectors = [
    SelectFromModel(Lasso(alpha=0.1, random_state=0)),
    'passthrough'
]

steps = [('scaler', StandardScaler()), # preprocessing steps

```

```

        ('feature_selector', SelectFromModel(Lasso(alpha=0.1,
random_state=0))), # Feature selection
        ('clf', RandomForestClassifier(random_state=0))]
# Model

param_grid = {
    'scaler': scalers,
    'feature_selector': feature_selectors,
    'clf__n_estimators': [10, 1000],
    'clf__max_depth': [1, None]
}

pipeline = Pipeline(steps)

search = GridSearchCV(pipeline, param_grid, n_jobs=-1, cv = 5, scoring
= "balanced_accuracy")
search.fit(X, y)
print("Best parameter (CV score=%0.2f):" % search.best_score_)
print(search.best_params_)

```

```

Best parameter (CV score=0.81):
{'clf__max_depth': None, 'clf__n_estimators': 1000,
'feature_selector': SelectFromModel(estimator=Lasso(alpha=0.1,
random_state=0)), 'scaler': StandardScaler()}

```

```

results = pd.DataFrame(search.cv_results_)
results.sort_values('rank_test_score')[[
    'param_clf__max_depth', 'param_clf__n_estimators',
    'param_feature_selector', 'param_scaler', 'params',
    'mean_test_score', 'std_test_score',
    'rank_test_score']]

```

	param_clf__max_depth	param_clf__n_estimators	\
12	None	1000	
14	None	1000	
13	None	1000	
15	None	1000	
8	None	10	
9	None	10	
10	None	10	
11	None	10	
0	1	10	
1	1	10	
2	1	10	
3	1	10	
4	1	1000	
5	1	1000	
6	1	1000	
7	1	1000	

```

param_feature_selector
param_scaler \
12 SelectFromModel(estimator=Lasso(alpha=0.1, ran...
StandardScaler()
14                                     passthrough
StandardScaler()
13 SelectFromModel(estimator=Lasso(alpha=0.1, ran...
passthrough
15                                     passthrough
passthrough
8   SelectFromModel(estimator=Lasso(alpha=0.1, ran...
StandardScaler()
9   SelectFromModel(estimator=Lasso(alpha=0.1, ran...
passthrough
10                                     passthrough
StandardScaler()
11                                     passthrough
passthrough
0   SelectFromModel(estimator=Lasso(alpha=0.1, ran...
StandardScaler()
1   SelectFromModel(estimator=Lasso(alpha=0.1, ran...
passthrough
2                                     passthrough
StandardScaler()
3                                     passthrough
passthrough
4   SelectFromModel(estimator=Lasso(alpha=0.1, ran...
StandardScaler()
5   SelectFromModel(estimator=Lasso(alpha=0.1, ran...
passthrough
6                                     passthrough
StandardScaler()
7                                     passthrough
passthrough

```

	params	mean_test_score
\		
12	{'clf__max_depth': None, 'clf__n_estimators': ...	0.806935
14	{'clf__max_depth': None, 'clf__n_estimators': ...	0.806935
13	{'clf__max_depth': None, 'clf__n_estimators': ...	0.801380
15	{'clf__max_depth': None, 'clf__n_estimators': ...	0.801380
8	{'clf__max_depth': None, 'clf__n_estimators': ...	0.770950
9	{'clf__max_depth': None, 'clf__n_estimators': ...	0.770950

10	{'clf__max_depth': None, 'clf__n_estimators': ...	0.770950
11	{'clf__max_depth': None, 'clf__n_estimators': ...	0.770950
0	{'clf__max_depth': 1, 'clf__n_estimators': 10,...	0.624183
1	{'clf__max_depth': 1, 'clf__n_estimators': 10,...	0.624183
2	{'clf__max_depth': 1, 'clf__n_estimators': 10,...	0.624183
3	{'clf__max_depth': 1, 'clf__n_estimators': 10,...	0.624183
4	{'clf__max_depth': 1, 'clf__n_estimators': 100...	0.594164
5	{'clf__max_depth': 1, 'clf__n_estimators': 100...	0.594164
6	{'clf__max_depth': 1, 'clf__n_estimators': 100...	0.594164
7	{'clf__max_depth': 1, 'clf__n_estimators': 100...	0.594164

	std_test_score	rank_test_score
12	0.045989	1
14	0.045989	1
13	0.043601	3
15	0.043601	3
8	0.049943	5
9	0.049943	5
10	0.049943	5
11	0.049943	5
0	0.040118	9
1	0.040118	9
2	0.040118	9
3	0.040118	9
4	0.055740	13
5	0.055740	13
6	0.055740	13
7	0.055740	13

```
results[['split0_test_score',
        'split1_test_score', 'split2_test_score', 'split3_test_score',
        'split4_test_score', 'mean_test_score',
        'std_test_score']].sort_values('std_test_score')
```

	split0_test_score	split1_test_score	split2_test_score	\
0	0.676471	0.583333	0.666667	
1	0.676471	0.583333	0.666667	
2	0.676471	0.583333	0.666667	
3	0.676471	0.583333	0.666667	

13	0.788646	0.809524	0.730159
15	0.788646	0.809524	0.730159
12	0.788646	0.837302	0.730159
14	0.788646	0.837302	0.730159
8	0.747606	0.718254	0.753968
9	0.747606	0.718254	0.753968
10	0.747606	0.718254	0.753968
11	0.747606	0.718254	0.753968
4	0.617647	0.527778	0.658730
5	0.617647	0.527778	0.658730
6	0.617647	0.527778	0.658730
7	0.617647	0.527778	0.658730

	split3_test_score	split4_test_score	mean_test_score
std_test_score			
0	0.583333	0.611111	0.624183
0.040118			
1	0.583333	0.611111	0.624183
0.040118			
2	0.583333	0.611111	0.624183
0.040118			
3	0.583333	0.611111	0.624183
0.040118			
13	0.813492	0.865079	0.801380
0.043601			
15	0.813492	0.865079	0.801380
0.043601			
12	0.813492	0.865079	0.806935
0.045989			
14	0.813492	0.865079	0.806935
0.045989			
8	0.769841	0.865079	0.770950
0.049943			
9	0.769841	0.865079	0.770950
0.049943			
10	0.769841	0.865079	0.770950
0.049943			
11	0.769841	0.865079	0.770950
0.049943			
4	0.527778	0.638889	0.594164
0.055740			
5	0.527778	0.638889	0.594164
0.055740			
6	0.527778	0.638889	0.594164
0.055740			
7	0.527778	0.638889	0.594164
0.055740			