

Lab 6 - Bayesian Knowledge Tracing (BKT) and Variants

This tutorial is partially based on the pyBKT model tutorial and the Jupyter notebooks available on GitHub at <https://github.com/CAHLR/pyBKT>.

One notable application of machine learning in education is represented **knowledge inference models**, which aim to understand how well a student is learning concepts or skills. Being able to monitor this knowledge makes it possible to improve and personalize online learning platforms or intelligent tutoring systems, by focusing on areas the student is weak in and accelerating learning of certain concepts.

In this tutorial, we study a range of popular models for modelling students' knowledge based on **Bayesian Knowledge Tracing (BKT)**. BKT was introduced in 1995 as a means to model students' knowledge as a **latent variable** in online learning environments. Specifically, the environment can maintain an estimate of the **probability that the student has learned a set of skills**, which is statistically equivalent to a 2-node dynamic Bayesian network.

For this tutorial, we will rely on a Python implementation of the Bayesian Knowledge Tracing algorithm and more recent variants, estimating student cognitive mastery from problem solving sequences, known under the name of **pyBKT**. This package can be used to define and fit many BKT variants.

These variants are derived from a range of papers published in the educational data mining literature and, in this tutorial, we will provide you with the main notions and implementation details needed to investigate BKT models in practice.

Expected Tasks

- Follow the pyBKT getting started showcase.
- Solve a range of exercises on BKT models.

Learning Objectives

- Instantiate and run a pipeline on BKT models.
- Conduct fine-grained analyses on specific learning skills.
- Understand and experiment with different variants of BKT.
- Compare the performance of BKT setups under different evaluation methods.
- Inspect the influence of a BKT variant on the internal BKT parameters.

More information on the PyBKT is provided in the corresponding [Github repository](#).

```
# Traditional packages
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
```

```
import math
```

```
%matplotlib inline
```

Introduction

BKT models operationalize the learning of a student as a **Markov process**, building upon the idea that, while students interact with an educational environment, their skill in a given concept improves. To move the theory behind BKT into practice, variables related to forgetting, learning, guessing, slipping, and so on need to be modelled, controlling for instance how fast and how well learning is happening for the student.

The BKT model assumes that the student's knowledge can be estimated by means of standardized questions, which can be answered correctly or incorrectly, on a concept or combination of concepts. BKT also assumes that initially a student may not know about a concept, but their knowledge gets better with learning and practice related to that concept. The following concepts will be

- P_0 is the initial probability of mastering that concept (skill).
- P_F is the probability that the student forgot something previously learned on the concept (skill).
- P_L is the probability that the student has learned something that was previous not known on the concept (skill).
- P_S is the probability that the student gave a wrong answer even though they had learned the concept (skill).
- P_G is the probability that the student guessed the right answer while not knowing the concept (skill).

In this tutorial, we will use a dataset of the student's responses to questions in a test, along with whether they answered correctly or incorrectly, and we will use a BKT model to find the values of the above probabilities.

The ASSISTments data set

ASSISTments is a free tool for assigning and assessing math problems and homework. Teachers can select and assign problem sets. Once they get an assignment, students can complete it at their own pace and with the help of hints, multiple chances, and immediate feedback. Teachers get instant results broken down by individual student or for the whole class. Please, find more information on the platform [here](#).

In this tutorial, we will play with a simplified version of a dataset collected from the ASSISTments tool, saved on a CSV files with the following columns:

- `user_id`: The ID of the student doing the problem.
- `template_id`: The ID of the template in ASSISTment (assistments with the same template ID have similar questions).

- `assistment_id`: The ID of the ASSISTment (an assistment consists of one or more problems).
- `order_id`: These IDs are chronological and refer to the id of the original problem log.
- `problem_id`: The ID of the problem.
- `skill_name`: Skill name associated with the problem.
- `correct`: 1 if correct on the first attempt, 0 if incorrect on the first attempt or asked for help.
- `ms_first_response`: The time in milliseconds for the student's first response.
- `attempt_count`: Number of student attempts on this problem.
- `hint_count`: Number of student hints asked by the student on this problem.
- `hint_total`: Number of possible hints to be asked on this problem.

```
DATA_DIR = "../../../data/"
as_data = pd.read_csv(DATA_DIR + 'as_supersmall.csv',
encoding='latin', low_memory=False)
```

```
as_data.head(10)
```

	user_id	template_id	assistment_id	order_id	problem_id	
0	70733	30060	33175	35278766	51460	Box and Whisker
1	70733	30060	33182	35278780	51467	Box and Whisker
2	70733	30059	33107	35278789	51392	Box and Whisker
3	70733	30060	33187	35278802	51472	Box and Whisker
4	70733	30059	33111	35278810	51396	Box and Whisker
5	70872	30059	33136	32268742	51421	Box and Whisker
6	70872	30799	33144	32268764	51429	Box and Whisker
7	72059	30799	33155	33409110	51440	Box and Whisker
8	72059	30060	33181	33409165	51466	Box and Whisker
9	72059	30060	33168	33409366	51453	Box and Whisker

	correct	ms_first_response	attempt_count	hint_count	hint_total
0	0	9575	2	0	4
1	1	6422	1	0	4
2	0	11365	3	0	3
3	1	4412	1	0	4
4	1	6902	1	0	3
5	1	7281	1	0	3
6	1	7234	1	0	3

7	0	38290	2	0	3
8	0	8366	4	0	4
9	1	9661	1	0	4

Before delving into the pyBKT description and showcase, we invite you to spend some time to explore the toy dataset presented in this tutorial, e.g., how many students/problems/skills are included, examine the skills in more detail etc. Here, you could therefore add one or more cells to perform your exploration.

The pyBKT Package

In this tutorial, we use the pyBKT package, a Python implementation of the Bayesian Knowledge Tracing algorithm and variants, estimating student cognitive mastery from problem solving sequences. First, we install the package:

```
%pip install pyBKT
```

```
Defaulting to user installation because normal site-packages is not
writeable
Requirement already satisfied: pyBKT in /usr/local/lib/python3.8/dist-
packages (1.4)
Requirement already satisfied: requests in
/usr/local/lib/python3.8/dist-packages (from pyBKT) (2.28.1)
Requirement already satisfied: sklearn in
/usr/local/lib/python3.8/dist-packages (from pyBKT) (0.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.8/dist-
packages (from pyBKT) (1.22.4)
Requirement already satisfied: pandas in
/usr/local/lib/python3.8/dist-packages (from pyBKT) (1.4.3)
Requirement already satisfied: python-dateutil>=2.8.1 in
/usr/local/lib/python3.8/dist-packages (from pandas->pyBKT) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.8/dist-packages (from pandas->pyBKT) (2022.2.1)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.8/dist-packages (from requests->pyBKT)
(2022.6.15)
Requirement already satisfied: charset-normalizer<3,>=2 in
/usr/local/lib/python3.8/dist-packages (from requests->pyBKT) (2.1.1)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in
/usr/local/lib/python3.8/dist-packages (from requests->pyBKT)
(1.26.12)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.8/dist-packages (from requests->pyBKT) (3.3)
Requirement already satisfied: scikit-learn in
/usr/local/lib/python3.8/dist-packages (from sklearn->pyBKT) (1.0.2)
Requirement already satisfied: six>=1.5 in /usr/lib/python3/dist-
packages (from python-dateutil>=2.8.1->pandas->pyBKT) (1.14.0)
Requirement already satisfied: joblib>=0.11 in
/usr/local/lib/python3.8/dist-packages (from scikit-learn->sklearn-
>pyBKT) (1.1.0)
```

Requirement already satisfied: scipy>=1.1.0 in
/usr/local/lib/python3.8/dist-packages (from scikit-learn->sklearn-
>pyBKT) (1.9.1)

Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.8/dist-packages (from scikit-learn->sklearn-
>pyBKT) (3.1.0)

Note: you may need to restart the kernel to use updated packages.

Then, we can import the core class provided by the package, that is Model.

```
from pyBKT.models import Model
```

The first step is to construct a BKT model. To be instantiated, a BKT model requires a series of parameters, whose default value and meaning is provided below (e.g., num_fits, seed, defaults, and any model variant(s) that may be used). Each parameter can be modified during fit/crossvalidation time too.

- **Defaults generic parameters:**

- num_fits (5) is the number of initialization fits used for the BKT model.
- defaults (None) is a dictionary that can be used to pass values different than the default ones during initialization.
- parallel (True) indicates whether the computation will use multi-threading.
- skills ('.*') is a regular expression used to indicate the skills the BKT model will be run on.
- seed (random.randint(0, 1e8)) is a seed that can be setup to enable reproducible experiments.
- folds (5) is the number of folds used in case of cross-validation.
- forgets (False) indicates whether the model will consider that the student may give a wrong answer even though they had learned the concept.

- **Defaults additional parameters:**

- order_id ('order_id') is the name of the CSV column for the chronological IDs that refer to the original problem log.
- skill_name ('skill_name') is the name of the CSV column for the skill name associated with the problem.
- correct ('correct') is the name of the CSV column for the correct / incorrect label on the first attempt.
- user_id ('user_id') is the name of the CSV column for the ID of the student doing the problem.
- multilearn ('template_id') is the name of the column for checking whether there is a multi-skill object.
- multiprior ('correct') is the name of the CSV column for mapping multi-prior knowledge.
- multigs ('template_id') is the name of the CSV column corresponding to the desired guess/slip classes.

- **Initializers for learnable parameters:**
 - 'prior' (None, no initialization) is the initial probability of answering the question correct.
 - 'learns' (None, no initialization) is the probability that the student has learned something that was previous not known.
 - 'guesses' (None, no initialization) is the probability that the student guessed the right answer while not knowing the concept.
 - 'slips' (None, no initialization) is the probability that the student gave a wrong answer even though they had learned the concept.
 - 'forgets' (None, no initialization) is the probability that the student forgot something previously learned.

If you have doubts on the meaning of certain parameters, please ask to TAs or move on the next examples (they will help you understand).

```
model = Model(seed=0)
model
```

```
Model(parallel=True, num_fits=5, seed=0, defaults=None)
```

The Model class is inspired by scikit-learn and, therefore, provides a range of methods a model can be called with:

- The **fit** method fits a BKT model given model and data information. Takes arguments skills, number of initialization fits, default column names (i.e. correct, skill_name), parallelization, and model types.
- The **predict** method predicts using the trained BKT model and test data information. Takes test data path or DataFrame as arguments. Returns a dictionary mapping skills to predicted values for those skills. Note that the predicted values are a tuple of (correct_predictions, state_predictions).
- The **evaluate** method evaluates a BKT model given model and data information. Takes a metric and data path or DataFrame as arguments. Returns the value of the metric for the given trained model tested on the given data.
- The **crossvalidate** method crossvalidates (trains and evaluates) the BKT model. Takes the data, metric, and any arguments that would be passed to the fit function (skills, number of initialization fits, default column names, parallelization, and model types) as arguments.

We will show a range of examples for each of the above methods.

Fitting and evaluating a model

```
model = Model(seed=0)
%time model.fit(data=as_data, skills='Box and Whisker')
%time model.evaluate(data=as_data, metric='auc')
```

```
CPU times: user 682 ms, sys: 12 ms, total: 694 ms
Wall time: 311 ms
```

```
CPU times: user 152 ms, sys: 0 ns, total: 152 ms
Wall time: 94.6 ms
```

```
0.5674965857758167
```

First, we have fitted a BKT model on the 'Box and Whisker' skill and, then, evaluate the corresponding **training AUC** (0.64). Note that we have run the BKT fitting process on the full dataset, to understand how well the BKT model can fit the data. Evaluation methods like cross-validation will be presented later in this notebook. Furthermore, the default metric displayed is RMSE, but pyBKT supports AUC ('auc'), RMSE ('rmse'), and accuracy ('accuracy') as metrics. We will also see how to add other metrics.

For each skill, you can get the learned parameters for 'prior', 'learns', 'guesses', 'slips', and 'forgets'. Specifically:

- **prior** (P_0): the prior probability of "knowing".
- **forgets** (P_F): the probability of transitioning to the "not knowing" state given "known".
- **learns** (P_L): the probability of transitioning to the "knowing" state given "not known".
- **slips** (P_S): the probability of picking incorrect answer, given "knowing" state.
- **guesses** (P_G): the probability of guessing correctly, given "not knowing" state.

```
model.coef_
```

```
{'Box and Whisker': {'prior': 0.8533977868556804,
  'learns': array([0.19638124]),
  'guesses': array([0.24682242]),
  'slips': array([0.21955686]),
  'forgets': array([0.])}}
```

```
model.params()
```

skill	param	class	value
Box and Whisker	prior	default	0.85340
	learns	default	0.19638
	guesses	default	0.24682
	slips	default	0.21956
	forgets	default	0.00000

We could initialize the prior knowledge to $1e-40$ for Box and Whisker, before fitting the model.

```
model = Model(seed=0)
```

```
model.coef_ = {'Box and Whisker': {'prior': 1e-40}}
model.coef_
```

```
{'Box and Whisker': {'prior': 1e-40}}
```

Then, we can fit the model and observe the resulting AUC score. How does it compares to the AUC score of the previous model.

```
%time model.fit(data=as_data, skills='Box and Whisker')
%time model.evaluate(data=as_data, metric='auc')
```

```
CPU times: user 261 ms, sys: 0 ns, total: 261 ms
Wall time: 130 ms
CPU times: user 112 ms, sys: 0 ns, total: 112 ms
Wall time: 48.8 ms
```

```
0.5535245298875933
```

```
model.params()
```

			value
skill	param	class	
Box and Whisker	prior	default	0.00000
	learns	default	0.56145
	guesses	default	0.62351
	slips	default	0.22136
	forgets	default	0.00000

You can also train simple BKT models on different skills in the data set.

```
model = Model(seed=0)
%time model.fit(data=as_data, skills=['Box and Whisker', 'Scatter
Plot'])
%time model.evaluate(data=as_data, metric='auc')
```

```
CPU times: user 248 ms, sys: 0 ns, total: 248 ms
Wall time: 149 ms
CPU times: user 86.6 ms, sys: 0 ns, total: 86.6 ms
Wall time: 54.3 ms
```

```
0.6625740160222918
```

And, then, observed the learned parameters for each skill. Note that, when multiple skills are passed to fit, the method will run a fitting procedure for each skill, separately (in this case, we will have two BKT models).

```
model.params()
```

			value
skill	param	class	
Box and Whisker	prior	default	0.96146
	learns	default	0.01810
	guesses	default	0.00036
	slips	default	0.23364
	forgets	default	0.00000
Scatter Plot	prior	default	0.48923
	learns	default	0.56128


```
guesses default 0.71140
slips    default 0.03956
forgets  default 0.00000
```

You can also enable forgetting, by setting the corresponding parameter in the fit method.

```
model = Model(seed=0)
%time model.fit(data=as_data, skills='Box and Whisker', forgets=True)
%time model.evaluate(data=as_data, metric='auc')
```

```
CPU times: user 73.4 ms, sys: 0 ns, total: 73.4 ms
```

```
Wall time: 49.4 ms
```

```
CPU times: user 125 ms, sys: 0 ns, total: 125 ms
```

```
Wall time: 39.8 ms
```

```
0.5721189200546275
```

```
model.params()
```

			value
skill	param	class	
Box and Whisker	prior	default	0.72657
	learns	default	0.25751
	guesses	default	0.36213
	slips	default	0.19171
	forgets	default	0.01321

Or train a multiguess and slip BKT model on the same skills in the data set. The **multigs** model fits a different guess/slip rate for each class. Note that, with *multigs=True*, the guess and slip classes will be specified by the *template_id*. You can specify a custom column mapping by doing *multigs='column_name'*.

```
model = Model(seed=0)
%time model.fit(data=as_data, skills=['Box and Whisker'],
multigs=True)
%time model.evaluate(data=as_data, metric='auc')
```

```
CPU times: user 173 ms, sys: 0 ns, total: 173 ms
```

```
Wall time: 62.1 ms
```

```
CPU times: user 128 ms, sys: 0 ns, total: 128 ms
```

```
Wall time: 116 ms
```

```
0.7049059775186469
```

And finally, we show the BKT paramaters. By enabling *multigs*, the guess and slip classes will be specified by the *template_id* and, by setting *multigs=True*, the guess and slip classes will be specified by default by the *template_id* classes. Note that assistments with the same template ID have similar questions. What could you observe by looking at the different learned guesses and slips values below?

```
model.params()
```

skill	param	class	value
Box and Whisker	prior	default	0.09283
	learns	default	0.07509
	guesses	30059	0.75371
		30060	0.59704
		30799	0.70751
		63446	0.00000
		63447	0.16724
		63448	1.00000
	slips	30059	0.01865
		30060	0.16663
		30799	0.04409
		63446	1.00000
		63447	0.99707
		63448	0.00000
	forgets	default	0.00000

The **multilearn** model fits a different learn rate (and forget rate if enabled) rate for each class specified. Note that, with multilearn=True, the learn classes are specified by the *template_id*. You can specify a custom column mapping by doing *multilearn='column_name'*.

```
model = Model(seed=0)
%time model.fit(data=as_data, skills=['Box and Whisker'],
multilearn=True)
%time model.evaluate(data=as_data, metric='auc')
```

```
CPU times: user 108 ms, sys: 0 ns, total: 108 ms
Wall time: 55.3 ms
CPU times: user 119 ms, sys: 3.37 ms, total: 122 ms
Wall time: 40.4 ms
```

```
0.5694925937598487
```

Looking at the parameters, we will observe a learns score for each template_id (the class column in the paras dataframe). In this case, what could you observe by looking at the different learns values below?

```
model.params()
```

skill	param	class	value
Box and Whisker	prior	default	0.94327
	learns	30059	0.02049
		30060	0.42985
		30799	0.25272
		63446	0.84344
		63447	0.66021
		63448	0.25677
	guesses	default	0.05763
	slips	default	0.23095

forgets	30059	0.00000
	30060	0.00000
	30799	0.00000
	63446	0.00000
	63447	0.00000
	63448	0.00000

You can also combine multiple variants, and use a different column to specify the different learn and forget classes. In this case, we use `user_id`, assuming that we are interested in learning the parameters for each student, and we also enable forgetting.

```
model = Model(seed=0)
%time model.fit(data=as_data, skills=['Box and Whisker'],
forgets=True, multilearn='user_id')
%time model.evaluate(data=as_data, metric='auc')
```

CPU times: user 777 ms, sys: 1.33 ms, total: 778 ms

Wall time: 368 ms

CPU times: user 106 ms, sys: 362 µs, total: 106 ms

Wall time: 121 ms

0.6405609832965647

Once we run a BKT model with `forgets=True` and `multilearn='user_id'`, we will observe individual scores for each student, as shown below.

```
model.params()
```

skill	param	class	value
Box and Whisker	prior	default	0.05051
		70733	0.08695
	learns	70872	0.11881
		72059	0.07787
		79748	0.11881
		79750	0.13948
		79769	0.10011
		81641	0.02619
		82482	0.12312
		82533	0.16158
		83169	0.11881
		84316	0.07571
		85725	0.13006
		86489	0.09430
		91260	0.11881
		91301	0.11881
		91406	0.16158
		91409	0.11659
		91436	0.01063
		91459	0.10580
		96217	0.10079

	96242	0.07787
	96248	0.06080
	96249	0.06670
	96266	0.11881
	96268	0.08741
	96292	0.07408
	96294	0.11825
	96296	0.11028
guesses	default	0.70234
slips	default	0.05860
forgets	70733	0.31890
	70872	0.23699
	72059	0.35471
	79748	0.23699
	79750	0.22281
	79769	0.26551
	81641	0.66799
	82482	0.26814
	82533	0.19007
	83169	0.23699
	84316	0.35968
	85725	0.23422
	86489	0.34166
	91260	0.23699
	91301	0.23699
	91406	0.19007
	91409	0.28950
	91436	0.78903
	91459	0.32115
	96217	0.27475
	96242	0.35471
	96248	0.40840
	96249	0.42812
	96266	0.23699
	96268	0.26268
	96292	0.36373
	96294	0.27470
	96296	0.26723

The best performing models are typically those that combine several useful variants, such as the multilearn and multiguess/slip class variants. After this lab session, you might be interested in testing with other skills and see whether this observations is true for other skills as well.

Make predictions

As we said, the predict method can be executed on the trained BKT model, obtaining a dictionary mapping skills to predicted values for those skills, namely correct_predictions (each score is between 0 and 1 that measures the extent to which the model thinks that the student will answer correctly to that question) and state_predictions (each score between 0

and 1 that measures the extent to which the student has mastered that skill, after that question).

Note that, in the example below, we have run the BKT fitting process on the full dataset, to understand how well the BKT model can fit the data. Evaluation methods like cross-validation will be presented slightly after in this notebook.

```
model = Model(seed=0)
%time model.fit(data=as_data, skills='Box and Whisker')
%time model.evaluate(data=as_data, metric='auc')
%time preds = model.predict(data=as_data)
```

```
CPU times: user 83.1 ms, sys: 4.64 ms, total: 87.7 ms
Wall time: 48.6 ms
CPU times: user 137 ms, sys: 0 ns, total: 137 ms
Wall time: 57.2 ms
CPU times: user 195 ms, sys: 0 ns, total: 195 ms
Wall time: 98 ms
```

```
preds[preds['skill_name']=='Box and Whisker'][['user_id', 'correct',
'correct_predictions', 'state_predictions']]
```

	user_id	correct	correct_predictions	state_predictions
0	70733	0	0.66780	0.81542
1	70733	1	0.59418	0.69635
2	70733	0	0.74813	0.94535
3	70733	1	0.70860	0.88142
4	70733	1	0.77085	0.98210
...
219	96296	1	0.78191	1.00000
220	96296	1	0.78191	1.00000
221	96296	1	0.78191	1.00000
222	96296	1	0.78191	1.00000
223	96296	1	0.78191	1.00000

```
[224 rows x 4 columns]
```

Note that, if the BKT model is asked to predict on skills not included in the training set, the output predictions for that skills will be a best effort guess of 0.5 for both the correct and state predictions.

Extend the evaluation

The pyBKT package makes also possible to extend the range of metrics you can compute while evaluating a BKT model. To this end, you need to define a custom function that, given true_vals (true values for the correct target) and pred_vals (the predicted values for the correct target), computes and returns the score corresponding to the desired metric.

```
def mae(true_vals, pred_vals):
    return np.mean(np.abs(true_vals - pred_vals))
```

```
%time model.evaluate(data=as_data, metric=mae)
```

```
CPU times: user 200 ms, sys: 1.41 ms, total: 201 ms
```

```
Wall time: 60.9 ms
```

```
0.3699436448390422
```

Perform cross validation

Finally, the pyBKT package offers also a cross-validation method. You can specify the number of folds, a seed, and a metric (one of the three default ones, namely 'rmse', 'auc' or 'accuracy', or a custom Python function as we have seen above). Furthermore, similarly to the fit method, arguments for cross-validation a BKT variant and for defining the data path/data and skill names are accepted.

```
model = Model(seed=0)
```

```
%time model.crossvalidate(data=as_data, skills='Box and Whisker',  
folds=5, metric='auc')
```

```
CPU times: user 3.06 s, sys: 3.31 ms, total: 3.06 s
```

```
Wall time: 1.5 s
```

```
          auc  
skill  
Box and Whisker 0.54551
```

In this showcase, we just opted to five folds due to the time constraints. In the other cases, you need to select an appropriate number of folds based on the data you are dealing with, as discussed in the lectures.

Exercises

That's your turn! We ask you to complete the following exercises. In case you do not finish them during the lab session, please feel free to complete later, at your earliest convenience. TAs are happy to address any question or doubts you might have.

Kindly note that the following exercises have the goal of supporting you in getting familiar with the library functions, and may not fully represent the sequences of steps and the design choices made in a real-world or homework scenario. Elements concerning the latter scenarios will be discussed during the session. Furthermore, due to running time constraints, the following exercises will be run in a train-test split or full data set mode, while we leave the opportunity to run them under a cross-validation setting after this lab session.

In all your models, we ask you to set the *seed* to 0, to let you reproduce the same results across different runs.

Note that the expected running time may vary according to the device or environment.

Question 1 [expected total time for BKT fitting: 2 mins]

- Fit a BKT model with default parameters on the full data set, only for the skill 'Addition and Subtraction Integers'.
- Compute the correct predictions from the BKT model, by using the predict method of the Model class.
- Manually calculate the RMSE between the true correct value and the predicted correct value (refer to Slide 51 of Lecture 4 to get the RMSE formula).
- Compare with the RMSE returned by the evaluate method of the BKT model.

EXERCISE CELL

```
def rmse(y_true, y_pred):  
    return np.sqrt(np.mean((y_true - y_pred) ** 2))  
  
model = Model(seed=0)  
%time model.fit(data=as_data, skills='Addition and Subtraction  
Integers')  
preds = model.predict(data=as_data)  
  
preds_filtered = preds[preds['skill_name'].str.contains('Addition and  
Subtraction Integers')]  
manual_training_rmse = rmse(preds_filtered['correct'],  
preds_filtered['correct_predictions'])  
  
print('Manual RMSE:', manual_training_rmse)  
print('pyBKT RMSE:', model.evaluate(data=as_data, metric='rmse'))  
  
CPU times: user 2.53 s, sys: 15.3 ms, total: 2.55 s  
Wall time: 1.28 s  
Manual RMSE: 0.46390773499967014  
pyBKT RMSE: 0.46390773499967064
```

Question 2 [expected total time for BKT fitting: 7 mins]

- Perform a user-based train-test split of the data, with 20% of the users in the test set.
- Fit the two BKT model variants on the training set, only for the skill 'Addition and Subtraction Integers'.
 - default;
 - forgets=True;
- Which model variant listed below has the highest test AUC for 'Addition and Subtraction Integers' in the test set?

EXERCISE CELL

```
users = as_data['user_id'].unique()  
users_train = list(np.random.choice(users, int(len(users) * 0.8),  
replace=False))  
users_test = list(set(users) - set(users_train))  
  
X_train, X_test = as_data[as_data['user_id'].isin(users_train)],  
as_data[as_data['user_id'].isin(users_test)]
```

```
models = {}

model = Model(seed=0, num_fits=1)
%time model.fit(data=X_train, skills='Addition and Subtraction
Integers')
models['simple'] = model.evaluate(data=X_test, metric='auc')

model = Model(seed=0, num_fits=1)
%time model.fit(data=X_train, skills='Addition and Subtraction
Integers', forgets=True)
models['forgets'] = model.evaluate(data=X_test, metric='auc')

df = pd.DataFrame(models.items())
df.columns = ['Model Type', 'AUC']
df.set_index('Model Type')

CPU times: user 447 ms, sys: 4.42 ms, total: 451 ms
Wall time: 182 ms
CPU times: user 733 ms, sys: 0 ns, total: 733 ms
Wall time: 388 ms
```

	AUC
Model Type	
simple	0.55297
forgets	0.53747

Question 3 [expected total time for BKT fitting: 3 mins]

- Bin values in the `ms_first_response` column in `as_data` to categories ('less than 10s', 'less than 20s', 'less than 30s', 'less than 40s', 'less than 50s', 'other').
- Fit BKT models with different learn rates, according to the `ms_first_response` categories above, on the full data set, only for the skill 'Addition and Subtraction Integers'. You need to play with the `multilearn` parameter of the BKT fit method.
- Create a bar plot to show the P_L (learns) value for each `ms_first_response` category above. You basically need to play with the dataframe returned by `model.params()`, to prepare the data to be shown in the plot.
- Does binned response time influence the P_L parameter for the skill 'Addition and Subtraction Integers'? Which bin result in the highest P_L scores?

EXERCISE CELL

```
skill = 'Addition and Subtraction Integers'
```

```
# Binning
```

```
learn_maps = {0: 'less than 10s', 1: 'less than 20s', 2: 'less than
30s', 3: 'less than 40s', 4: 'less than 50s'}
as_data['bin_s_first_response'] = (as_data['ms_first_response'] // (10
* 1000)).map(learn_maps).fillna('other')
```



```
# Modelling
```

```
model = Model(seed=0, num_fits=1)
%time model.fit(data=as_data, skills=skill,
multilearn='bin_s_first_response')
params = model.params().reset_index()
```


```
params
```

```
CPU times: user 496 ms, sys: 0 ns, total: 496 ms
```

```
Wall time: 272 ms
```

	skill	param	class	value
0	Addition and Subtraction Integers	prior	default	0.74974
1	Addition and Subtraction Integers	learns	less than 10s	0.01536
2	Addition and Subtraction Integers	learns	less than 20s	0.02294
3	Addition and Subtraction Integers	learns	less than 30s	0.01146
4	Addition and Subtraction Integers	learns	less than 40s	0.01898
5	Addition and Subtraction Integers	learns	less than 50s	0.02503
6	Addition and Subtraction Integers	learns	other	0.01188
7	Addition and Subtraction Integers	guesses	default	0.37005
8	Addition and Subtraction Integers	slips	default	0.20730
9	Addition and Subtraction Integers	forgets	less than 10s	0.00000
10	Addition and Subtraction Integers	forgets	less than 20s	0.00000
11	Addition and Subtraction Integers	forgets	less than 30s	0.00000
12	Addition and Subtraction Integers	forgets	less than 40s	0.00000
13	Addition and Subtraction Integers	forgets	less than 50s	0.00000
14	Addition and Subtraction Integers	forgets	other	0.00000

```
plt.figure(figsize = (12, 6))
plt.title(r'$P_{L}$ per response time')
params_learns = params[params['param'] ==
'learns'].sort_values(by='class').copy()
labels = params_learns['class']
values = params_learns['value']
plt.bar(labels, values, width=0.4)
plt.ylabel(r'$P_{L}$')
plt.xlabel('Response Time')
plt.ylim([0, .1])
plt.grid(axis='y')
plt.show()
```



Question 4 [expected total time for BKT fitting: 8 mins]

- Use the same bins `ms_first_response` to categories ('less than 10s', 'less than 20s', 'less than 30s', 'less than 40s', 'less than 50s', 'other').
- Fit a BKT model with `template-id` multilearn (default), on the full data set, only for the skill 'Addition and Subtraction Integers'.
- Fit a BKT model with `binned-response-time-based` multilearn, on the full data set, only for the skill 'Addition and Subtraction Integers'.
- Does the `binned-response-time-based` multilearn improve the AUC of the model compared to the default `template_id`-based multilearn?

EXERCISE CELL

```
skill = 'Addition and Subtraction Integers'
```

```
model = Model(seed=0, num_fits=1)
%time model.fit(data=as_data, skills=skill, multilearn=True)
default_multilearn_auc = model.evaluate(data=as_data, metric='auc')
```

```
model = Model(seed=0, num_fits=1)
%time model.fit(data=as_data, skills=skill,
multilearn='bin_s_first_response')
time_multilearn_auc = model.evaluate(data=as_data, metric='auc')
```

```
'AUC Improvement using Response Time:', time_multilearn_auc -
default_multilearn_auc
```

```
CPU times: user 455 ms, sys: 17.5 ms, total: 473 ms
Wall time: 162 ms
CPU times: user 420 ms, sys: 8.13 ms, total: 428 ms
Wall time: 205 ms
```

```
('AUC Improvement using Response Time:', -0.0031074282332050895)
```

Summary

In this tutorial, we have seen several important aspects of Bayesian Knowledge Tracing (BKT). We have shown how a typical data set for knowledge tracing should look like. We have illustrated how BKT models can be trained on different skills. We have shown how different variants of BKT can help you improve the goodness of your model. Many of the ideas described in this tutorial can be adapted to other data sets and projects. Finally, we have shown some examples of predictions and evaluations, covering also cross-validation. If you are interested in the implementation details of the different variants, we invite you to explore the codebase stored in the pyBKT Github repository.